

Assignment 5

Due November 25, 17:00

NOTE: Late submissions will **NOT** be accepted. Please submit a single PDF file with your answers via the **ECE 355 Brightspace** webpage.

1. [5 points] Consider the code portion of the matrix-vector product computation as shown below: (double) **128x128** matrix **A** is multiplied by (double) **128x1** vector **X**, producing (double) **128x1** result **Y** (initially all 0's).

```
for (i = 0; i < 128; i++) {
    for (j = 0; j < 128; j++) {
        Y[i] = Y[i] + A[i][j]*X[j];
    }
}
```

Storing **X**, **Y**, and **A** (each double array element is 8 bytes in size) requires $128 \times 8 + 128 \times 8 + 128 \times 128 \times 8 = \mathbf{130KB}$ of memory. If the cache (assume fully associative) is smaller than 130KB, the above code will cause many misses, considerably slowing down the program execution. Alternatively, one can perform blocked computation: partition **A** into smaller blocks and perform the product computation block-by-block. If our data blocks can fit into the cache, such blocked computation may significantly outperform the original code.

Rewrite the code fragment above using blocked computation and letting matrix **A**'s blocks be of size **64x64** (i.e., 4 blocks total). Assuming that such blocking yields the best performance, what can you say about the size of the cache?

2. [10 points] Consider a C code fragment below, working on a given square matrix float **X[N][N]** (stored row by row, i.e., in the row-major order), where **N = 256**:

```
float trace = 0;
for (i = 0; i < N; i++) {
    trace = trace + X[i][i];           /* sum diagonal elements of X */
}
for (i = 0; i < N; i++) {
    for (j = 0; j < N; j++) {
        X[i][j] = X[i][j]/trace;      /* normalize elements of X */
    }
}
```

Determine the x-related page fault rate in the following two cases: 1) the main memory uses **1-KB** paging with four pages allocated for **x**, and 2) the main memory uses **4-KB** paging with only one page allocated for **x**. Initially, no part of **x** is in the main memory. **Note:** Type float is 32 bits (4 bytes).

3. [10 points] Consider a C code fragment below, working on a given positive matrix `float x[M][N]` (stored row by row, i.e., in the row-major order), where M = 128 and N = 512:

```
for (i = 0; i < M; i++) {  
    float max = 0;  
    for (j = 0; j < N; j++) {  
        float y = X[i][j];  
        if (max < y) max = y;    /* find max. element of row i */  
    }  
    for (j = 0; j < N; j++) {  
        X[i][j] = X[i][j]/max; /* normalize elements of row i */  
    }  
}
```

Determine the x-related page fault rate in the following two cases: 1) the main memory uses **1-KB** paging with four pages allocated for x, and 2) the main memory uses **4-KB** paging with only one page allocated for x. Initially, no part of x is in the main memory. **Note:** Type `float` is 32 bits (4 bytes).