

Assignment 5

Due November 22, 14:59

NOTE: Late submissions will **NOT** be accepted. Please put your solutions in the ECE 355 drop-box (ELW, second floor) – they will be collected at **15:00**.

1. [5 points] Consider the code portion of the matrix-vector product computation as shown below: (double) **128x128** matrix **A** is multiplied by (double) **128x1** vector **X**, producing (double) **128x1** result **Y** (initially all 0's).

```
for (i = 0; i < 128; i++) {
    for (j = 0; j < 128; j++) {
        Y[i] = Y[i] + A[i][j]*X[j];
    }
}
```

Storing **X**, **Y**, and **A** (each double array element is 8 bytes in size) requires $128*8 + 128*8 + 128*128*8 = \mathbf{130KB}$ of memory. If the cache (assume fully associative) is smaller than 130KB, the above code will cause many misses, considerably slowing down the program execution. Alternatively, one can perform blocked computation: partition **A** into smaller blocks and perform the product computation block-by-block. If our data blocks can fit into the cache, such blocked computation may significantly outperform the original code.

Rewrite the code fragment above using blocked computation and letting matrix **A**'s blocks be of size **64x64** (i.e., 4 blocks total). Assuming that such blocking yields the best performance, what can you say about the size of the cache?

2. [10 points] Consider a C code fragment below, working on a given square matrix double **X[N][N]** (stored row by row, i.e., in the row-major order), where **N = 256**:

```
for (i = 0; i < N; i++) {
    double sum = 0;
    for (j = 0; j < N; j++) {
        sum = sum + X[i][j];
    }
    for (j = 0; j < N; j++) {
        X[i][j] = X[i][j]/sum;
    }
}
```

Determine the x-related page fault rate in the following two cases: 1) the main memory uses **2-KB** paging with two pages allocated for **x**, and 2) the main memory uses **4-KB** paging with only one page allocated for **x**. Initially, no part of **x** is in the main memory. **Note:** Type double is 64 bits (8 bytes).

3. [10 points] Consider a C code fragment below, modifying a given square matrix `int x[N][N]` (stored row by row, i.e., in the row-major order), where `N = 256`:

```
int average, dev;
for (i = 0; i < N; i++) {
    average = 0;
    for (j = 0; j < N; j++) {
        average = average + X[i][j]; /* sum row elements */
    }
    average = average/N;              /* row average */
    for (j = 0; j < N; j++) {
        dev = X[i][j] - average;    /* deviation from average */
        X[i][j] = dev*dev;          /* deviation squared */
    }
}
```

Determine the `x`-related page fault rate in the following two cases: 1) the main memory uses **1-KB** paging with four pages allocated for `x`, and 2) the main memory uses **4-KB** paging with only one page allocated for `x`. Initially, no part of `x` is in the main memory. **Note:** Type `int` is 32 bits (4 bytes).