

Assignment 6

ECE 355

V00984826

Problem 1

Answer the following questions about IEEE-754 32-bit floating-point representation:

1. (10 points)

- (a) Convert the IEEE-754 number 1 11111111 000000000000000000000000 to its decimal format.
- (b) Convert the IEEE-754 number 0 00000000 110000000000000000000000 to its decimal format.
- (c) Represent the decimal number -0.625 in the 32-bit IEEE-754 format.
- (d) Given the two IEEE-754 numbers X and Y below:

$$X = 0\ 01111011\ 100000000000000000000000, \quad Y = 1\ 01111110\ 110100000000000000000000,$$

calculate $Z = X - Y$ in binary format. Then, convert the resulting Z from IEEE-754 binary format to its decimal equivalent.

Solution 1

Part a)

For 1 11111111 000000000000000000000000:

$$M_1 = 0$$

$$E = 255$$

Result = Not a number (NaN)

Part b)

For 0 00000000 110000000000000000000000:

$$S = 0$$

$$E = 0$$

$$M = 1.1 \rightarrow 0$$

$$E = 0$$

$$\begin{aligned} \text{Number} &= +1 \times 2^{-126} \times 0.11 \\ &= 0.75 \times 2^{-126} \\ &= 8.8162076 \times 10^{-39} \end{aligned}$$

Part c)

For -0.625 :

$$\begin{aligned}0.625 &= 0.5 + 0.125 \\&= 2^{-1} + 2^{-3} \\&= (0.101)_2 \times 2^0 \\&= (1.01)_2 \times 2^{-1} \\ \text{Sign} &= 1 \\ \text{Exponent} &= -1 + 127 = 126 = (01111110)_2 \\ \text{Significand} &= 0100000000000000000000 \\ \text{IEEE-754} &= 1\ 01111110\ 0100000000000000000000\end{aligned}$$

Part d)

Given two 32-bit IEEE-754 numbers:

$$\begin{aligned}X &= 0\ 01111011\ 1000000000000000000000 \\Y &= 1\ 01111110\ 1101000000000000000000\end{aligned}$$

For X:

$$X = 1.5 \times 2^{-4} = 0.09375$$

For Y:

$$Y = -1.8125 \times 2^{-1} = -0.90625$$

Therefore:

$$Z = X - Y = 0.09375 - (-0.90625) = 1.0$$

Converting to IEEE-754:

$$Z = 0\ 01111111\ 0000000000000000000000$$

Problem 2

(5 points) Consider a pipelined datapath consisting of five stages:

- **F** – Fetch the instruction from the memory,
- **D** – Decode the instruction and read the source register(s),
- **C** – Execute the ALU operation specified by the instruction,
- **M** – Execute the memory operation specified by the instruction,
- **W** – Write the result in the destination register.

Identify data hazards in the code below and insert **NOP** instructions where necessary.

```

1  ADD #4, R0, R0      ; R0 = R0 + 4
2  ADD #4, R2, R2      ; R2 = R2 + 4
3  MOV (R0), R1        ; R1 = MEMORY[R0]
4  MOV (R2), R3        ; R3 = MEMORY[R2]
5  SUB R2, R0, R4      ; R4 = R2 - R0
6  SUB R3, R1, R5      ; R5 = R3 - R1
7  MOV R4, (R2)        ; MEMORY[R2] = R4
8  MOV R5, (R0)        ; MEMORY[R0] = R5
9  ADD #4, R0, R0      ; R0 = R0 + 4
10 ADD #4, R2, R2      ; R2 = R2 + 4

```

Solution 2

To resolve the data hazards in the given code, **NOP** (no-operation) instructions are inserted where necessary. The modified code with **NOP** instructions is as follows:

```

1  ADD #4, R0, R0      ; R0 = R0 + 4
2  ADD #4, R2, R2      ; R2 = R2 + 4
3  NOP
4  NOP
5  MOV (R0), R1        ; R1 = MEMORY[R0]
6  MOV (R2), R3        ; R3 = MEMORY[R2]
7  SUB R2, R0, R4      ; R4 = R2 - R0
8  NOP
9  NOP
10 SUB R3, R1, R5      ; R5 = R3 - R1
11 MOV R4, (R2)        ; MEMORY[R2] = R4
12 NOP
13 NOP
14 MOV R5, (R0)        ; MEMORY[R0] = R5
15 NOP
16 NOP
17 NOP
18 ADD #4, R0, R0      ; R0 = R0 + 4
19 ADD #4, R2, R2      ; R2 = R2 + 4

```

Explanation:

- After the first two **ADD** instructions, **NOP** instructions are inserted to prevent data hazards caused by dependent instructions that use R0 and R2.
- For the **SUB** instructions, additional **NOP** instructions are added because the values of R0 and R2 (updated by **ADD** or **MOV**) are required before proceeding.
- Similarly, after the **MOV** instructions that store the results in memory, **NOP** instructions are added to allow time for memory updates.
- Additional **NOP** instructions are inserted before the final **ADD** instructions to ensure no unresolved hazards.

Problems 3 and 4

3. (2 points) Solve Problem 12.8 from the textbook.

4. (8 points) Solve Problem 12.7 from the textbook.

Hint: Declare a shared counter variable, e.g., `volatile int thread_id_counter`, initialize it to 0 in `main()`, and poll it by each thread as follows:

```
while (thread_id_counter != my_id);
```

Each thread must increment `thread_id_counter` after updating `global dot_product`.

Solution for Problem 3

Given $P = 8$ and Speedup = 5, we need to solve the equation:

$$5 = \frac{1}{1 - f + \frac{f}{8}}$$

Rewriting the equation:

$$1 - f + \frac{f}{8} = \frac{1}{5}$$

$$1 - f + \frac{f}{8} = 0.2$$

$$1 - 0.2 = f - \frac{f}{8}$$

$$0.8 = f \left(1 - \frac{1}{8}\right)$$

$$0.8 = f \cdot \frac{7}{8}$$

$$f = \frac{0.8 \cdot 8}{7} = 0.914$$

Thus, $f = 0.914$, meaning the application program must be 91% parallelizable.

Solution for Problem 4

Below is the code for solving the problem using threads and a shared counter for synchronization. The shared variable ensures that each thread updates the global `dot_product` in a coordinated manner and avoiding race conditions.

```
1 #include <stdio.h>
2 #include "threads.h"
3
4 #define N 100
5 #define P 4
6
7 double a[N], b[N];
8 double dot_product = 0.0; // Global variable for storing the result
9 volatile int thread_id_counter = 0; // Shared counter for
    synchronization
```

```

10 Barrier bar;
11
12 void ParallelFunction (void)
13 {
14     int my_id, i, start, end;
15     double s;
16     my_id = get_my_thread_id();
17     start = (N/P) * my_id;
18     end = (N/P) * (my_id + 1) - 1;
19
20     // Step 1: Compute partial dot product for the assigned range
21     s = 0.0;
22     for (i = start; i <= end; i++)
23         s = s + a[i] * b[i];
24
25     // Step 2: Synchronize threads to update global sum one at a time
26     while (thread_id_counter != my_id); // Wait for the current thread's
        turn
27     dot_product = dot_product + s;      // Update the global dot product
28     thread_id_counter = (thread_id_counter + 1) % P; // Allow the next
        thread access
29
30     // Step 3: Synchronize all threads using a barrier
31     barrier(&bar, P);
32 }
33
34 void main(void)
35 {
36     int i;
37
38     /* Initialize vectors a[] and b[] (details omitted for brevity). */
39     init_barrier(&bar);                // Initialize the barrier for
        thread synchronization
40     thread_id_counter = 0;              // Initialize the shared counter
41     dot_product = 0.0;                 // Initialize the global dot
        product
42
43     // Step 4: Create threads for parallel computation
44     for (i = 1; i < P; i++)
45         create_thread(ParallelFunction);
46     ParallelFunction();                // Execute the function in the
        main thread
47
48     // Print the final result
49     printf("The dot product is %g\n", dot_product);
50 }

```

Explanation:

- **Global variables:** The `dot_product` variable is used to store the final result of the dot product calculation. The `thread_id_counter` ensures that threads update the global variable sequentially, avoiding race conditions.

- **Synchronization:** The `while` loop synchronizes threads by making them wait for their turn to update the global `dot_product`. The modulo operation ensures that access rotates among the threads.
- **Barrier:** After updating the global variable, threads wait at a barrier to ensure all threads complete their computations before proceeding.
- **Thread creation:** Threads are created in a loop, and each thread computes a portion of the dot product based on its assigned range.
- **Main thread:** The main thread also participates in the computation by executing `ParallelFunction()`.

This implementation ensures correct and efficient parallel computation of the dot product.